

A decorative graphic on the left side of the slide. It consists of several vertical bars of varying heights and widths in shades of gray. To the right of these bars are several circles of different sizes, also in shades of gray, arranged in a cluster.

SQL

Database Systems (H)
Dr Chris Anagnostopoulos



ROADMAP

Ta Dah!

- **Structured Query Language (SQL);**
 - Create a database *schema* & *relations* in SQL;
 - Assign key/integrity/referential *constraints* in SQL;
- **SELECT clause for *selection* queries;**
 - *Multi-sets* and *Sets* in SQL
 - Dealing with NULL values
- **Nested Correlated & Uncorrelated Queries**

PHILOSOPHY OF THE DECLARATIVE LANGUAGE

- Structured Query Language by R. Boyce (*known* from BCNF) in **1974**.
- SQL is a **declarative** language, i.e.,
 - declare *what to do* rather than *how to do it*
 - different from procedural languages, e.g., Java, C, C++.
- First *official standard*: **SQL-92**
- Latest release: **SQL:2016**...
- **Advice:** *follow* standard SQL to be compliant with *most* of the Database Systems 🏖️

SQL: DATABASE SCHEMA



- Statement: `CREATE SCHEMA`

```
CREATE SCHEMA Company;
```

Each statement in SQL *ends* with a semicolon ‘;’

SQL: CREATE TABLE



- A *new* relation (*table* in SQL):
 - Specify the *name* of the relation
 - Specify *attributes*, their *types* (domain), *constraints*

```
CREATE SCHEMA Company;  
CREATE TABLE EMPLOYEE ...;
```

SQL: ATTRIBUTES & DOMAINS

- **Numeric data types**

- Integer numbers: **INT**
- Floating-point (real) numbers: **REAL** or **DECIMAL(*n*, *m*)**
- DECIMAL(**3**,**2**) has 3 digits; 2 digits after the decimal '.' e.g., **9.99**

- **Character/String data types**

- Fixed length: **CHAR(*n*)**
 - i.e., exactly *n* characters
 - e.g., CHAR(**5**) has exactly **5** characters like 'Chris'
- *Variable* length: **VARCHAR(*n*)**
 - i.e., from 0 to *n* characters
 - e.g., VARCHAR(**5**) has up to **5** characters like 'C', or, 'Ch', or 'Chris'

SQL: ATTRIBUTES & DOMAINS

- **Bit-string data types** (*sequence of bits: e.g., 0101100*)
 - Fixed length: **BIT(*n*)**
 - Varying length: **BIT VARYING(*n*)**
- **Boolean data type**
 - Values of `TRUE` or `FALSE` or `NULL`
 - SQL is a 3-valued *logic*...(*yes, no, and maybe*)
- **DATE data type**
 - Ten positions for YEAR, MONTH, and DAY in the form
YYYY-MM-DD
- **More, like `TIMESTAMP`, `DATE INTERVALS`, ...**
Visit: <https://www.postgresql.org/docs/9.5/static/datatype.html>

SQL: CREATE TABLE

CREATE TABLE EMPLOYEE					
<i>attributes</i>	(Fname	VARCHAR(15)	<i>domain (type)</i>	NOT NULL,	<i>constraints</i>
	Minit	CHAR,		NOT NULL,	
	Lname	VARCHAR(15)		NOT NULL,	
	Ssn	CHAR(9)		NOT NULL,	
	Bdate	DATE,			
	Address	VARCHAR(30),			
	Sex	CHAR,			
	Salary	DECIMAL(10,2),			
	Super_ssn	CHAR(9),			
	Dno	INT		NOT NULL,	
	PRIMARY KEY (Ssn),				
CREATE TABLE DEPARTMENT					
	(Dname	VARCHAR(15)		NOT NULL,	
	Dnumber	INT		NOT NULL,	
	Mgr_ssn	CHAR(9)		NOT NULL,	
	Mgr_start_date	DATE,			
	PRIMARY KEY (Dnumber),				
	UNIQUE (Dname).				
	FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn));				
CREATE TABLE DEPT_LOCATIONS					
	(Dnumber	INT		NOT NULL,	
	Dlocation	VARCHAR(15)		NOT NULL,	
	PRIMARY KEY (Dnumber, Dlocation),				
	FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber));				

SQL: VALUE CONSTRAINTS

- *Default* value of an attribute
 - **DEFAULT** {value}
 - **NULL** is not permitted for a attribute (**NOT NULL**)
 - e.g., **DNO INT NOT NULL DEFAULT 1;**
- **CHECK** clause (*range domain* constraint)
 - e.g., **Dnumber INT NOT NULL CHECK(Dnumber > 0 AND Dnumber < 21);**

SQL: KEY CONSTRAINTS

- **Key constraint:** a **primary key value** is unique (no duplicates);
- **Entity Integrity constraint:** a primary key cannot be NULL;
- Primary Key Clause:
 - Dnumber **INT NOT NULL, PRIMARY KEY** (Dnumber);
- **UNIQUE** clause, specifies *candidate* keys
 - Dname **VARCHAR(15) NOT NULL, UNIQUE** (Dname);

```
CREATE TABLE DEPARTMENT
( Dname          VARCHAR(15)          NOT NULL,
  Dnumber        INT                  NOT NULL,
  Mgr_ssn        CHAR(9)              NOT NULL,
  Mgr_start_date DATE,
  PRIMARY KEY (Dnumber),
  UNIQUE (Dname),
  FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );
```

SQL: REFERENTIAL CONSTRAINTS

- FOREIGN KEY clause in EMPLOYEE
FOREIGN KEY (Super_ssn) REFERENCES Employee(Ssn)
- FOREIGN KEY clause in DEPARTMENT
FOREIGN KEY (Mgr_ssn) REFERENCES Employee(Ssn)
- Triggered actions for **Mgr_ssn, Super_ssn** when **Ssn** is updated or deleted:
 - Action: **ON DELETE SET NULL/ DEFAULT/ CASCADE**
 - Action: **ON UPDATE SET NULL/ DEFAULT /CASCADE**
- **CASCADE** option *propagates* DELETE / UPDATE to *all* referential tuples!
 - e.g., when the primary key **Ssn** is updated, then all foreign keys *refer* to it should be updated: **ON UPDATE CASCADE**
 - e.g., when the primary key **Ssn** is deleted, then all foreign keys *refer* to this tuple should be deleted: **ON DELETE CASCADE**

IN-CLASS QUIZ



Q1: What is happening when we *delete* a department?

Q2: What is happening when we *delete* an employee?

CREATE TABLE EMPLOYEE

```
( ... ,  
  Dno          INT          NOT NULL      DEFAULT 1,  
  CONSTRAINT EMPPK  
    PRIMARY KEY (Ssn),  
  CONSTRAINT EMPSUPERFK  
    FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)  
      ON DELETE SET NULL      ON UPDATE CASCADE,  
  CONSTRAINT EMPDEPTFK  
    FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber)  
      ON DELETE SET DEFAULT   ON UPDATE CASCADE);
```

CREATE TABLE DEPARTMENT

```
( ... ,  
  Mgr_ssn CHAR(9)          NOT NULL      DEFAULT '888665555',  
  ... ,  
  CONSTRAINT DEPTPK  
    PRIMARY KEY (Dnumber),  
  CONSTRAINT DEPTSK  
    UNIQUE (Dname),  
  CONSTRAINT DEPTMGRFK  
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)  
      ON DELETE SET DEFAULT   ON UPDATE CASCADE);
```

CREATE TABLE DEPT_LOCATIONS

```
( ... ,  
  PRIMARY KEY (Dnumber, Dlocation),  
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)  
    ON DELETE CASCADE      ON UPDATE CASCADE);
```



SELECT-FROM-WHERE

```
SELECT    <attribute list>
FROM      <table list>
WHERE     <condition>;
```

- Declare *what* to retrieve, i.e., which are the **attributes of interest**
- Declare *from* where to retrieve, i.e., which is the **table/relation**
- Declare with what *condition* to retrieve, i.e., which are the **conditions**

But, not saying how to *implement* this, e.g.,

- *how* to load the data from disk to memory,
- *how* to *search* and *check* if a tuple satisfies the condition, etc.

SELECT-FROM-WHERE

Query 0: Which are the addresses of employees working in the department 4 *or* their salary is less than 31000.

SELECT Address
FROM EMPLOYEE
WHERE DNO = 4 **OR** Salary < 31000

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

SELECT-FROM-WHERE: JOIN & SELECT

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname='Research' AND Dnumber=Dno;

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

Fname	Lname	Address
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

Query 2. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.



Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber AND Mgr_ssn=Ssn AND
 Plocation='Stafford';

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

<u>Pnumber</u>	<u>Dnum</u>	<u>Lname</u>	<u>Address</u>	<u>Bdate</u>
10	4	Wallace	291Berry, Bellaire, TX	1941-06-20
30	4	Wallace	291Berry, Bellaire, TX	1941-06-20



TABLE AS A VARIABLE

//in Java.

```
int e = 5;
```

```
int s = 7;
```

//in SQL

```
EMPLOYEE AS e
```

```
EMPLOYEE AS s
```

- A relation might play different *roles* within a query, e.g., **employee** *might* be a *supervisee* and **employee** *might* be a *supervisor*...(used in recursive references...)

SELECT ...

FROM EMPLOYEE AS **E, EMPLOYEE AS **S****

WHERE...

Query 3. For each *employee*, retrieve the employee's first and last name and the first and last name of his or her *immediate* supervisor.



```
SELECT E.Fname, E.Lname, S.Fname, S.Lname
FROM   EMPLOYEE AS E, EMPLOYEE AS S
WHERE  E.Super_ssn=S.Ssn;
```

EMPLOYEE **E**

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

→ {John Smith, Franklin Wong}

EMPLOYEE **S**

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

IN-CLASS ACTIVITY [A1]



Query 4: For each *employee*, if their *supervisor* is a *manager* of a *department*, show the department name and department number.

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

IN-CLASS ACTIVITY [A1]



EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

- John is supervised by **33344555**, who is **Franklin**. Franklin is manager of department **Research** (No. **5**).
- Franklin** is supervised by **888665555**, who is **James**. James is manager of **HQ** (No. **1**)
- ...

```
SELECT D.Dname, D.Dnumber
FROM   EMPLOYEE AS E, EMPLOYEE AS S, DEPARTMENT AS D
WHERE  E.Super_ssn=S.Ssn AND S.Ssn = D.Mgr_Ssn;
```



R. Cartesius;1648

IF WHERE IS MISSING...

- Missing WHERE: *no condition* on tuple selection
- If FROM involves *two or more* relations, **avoid**; *unreasonable* tuples.
- **Why?** CROSS (*Cartesian*) PRODUCT: *all* possible tuple combinations!

```
SELECT  Ssn  
FROM    EMPLOYEE;
```

```
SELECT  Ssn, Dname  
FROM    EMPLOYEE, DEPARTMENT;
```

Each *tuple* from **EMPLOYEE** is *concatenated* with *each* tuple from **DEPARTMENT**...disaster, computationally heavy, and meaningless!

MISSING WHERE IS CATASTROPHE



R. Cartesius;1648

```
SELECT Ssn, Dname
FROM EMPLOYEE, DEPARTMENT
```

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

John...Research

John...Administration

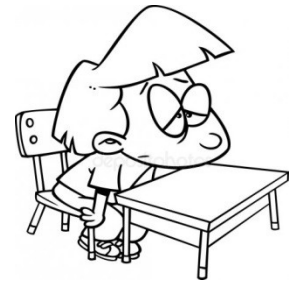
John...HQ

Franklin...Research

Franklin...Administration

...

USE OF THE ASTERISK



If *bored* listing *all* the attributes, then *use* asterisk (*), i.e., *all* attributes are of interest 📌

```
SELECT *  
FROM EMPLOYEE  
WHERE Dno=5;
```

Select all the information about **those** employees working at the department 5

```
SELECT *  
FROM EMPLOYEE, DEPARTMENT  
WHERE Dname='Research' AND Dno=Dnumber;
```

Select all the information (employee and department) from **those** employees working at the department 'Research'

```
SELECT *  
FROM EMPLOYEE, DEPARTMENT;
```

Select all the information about employees and departments **with no meaning** 📌

TABLES AS MULTI-SETS IN SQL

- **Set:** has only unique elements, e.g., $S = \{a, b, c\}$
- **Multiset:** might have duplicates, e.g., $M = \{a, a, a, b, c, c\}$
- Operators: UNION, EXCEPT, INTERSECT

Query 5: Retrieve the salary of *each* employee, and *all* the distinct salaries

SELECT
FROM

Salary
EMPLOYEE;

SELECT
FROM

DISTINCT Salary
EMPLOYEE;

Salary

10000

10000

25000

30000

25000

30000

30000

Salary

10000

25000

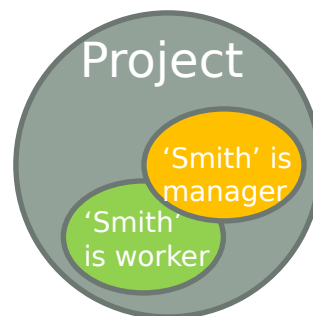
30000



IN-CLASS ACTIVITY [A2]

List *all* project numbers for projects that involve employees, whose last name is 'Smith', *either* as **workers** or as **managers** of departments controlling these projects.

Idea: *split* this into *two* sub-queries and then use the set UNION operator over the partial results.





IN-CLASS ACTIVITY [A2]

Step 1: Retrieve the projects where an employee with surname 'Smith' is *working on*;

Associate **EMPLOYEE** with **PROJECT** via **WORKS_ON**

```
( SELECT      DISTINCT Pnumber
  FROM        PROJECT, WORKS_ON, EMPLOYEE
 WHERE        Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```



IN-CLASS ACTIVITY [A2]

Step 2: Retrieve the projects where an employee with surname 'Smith' is a *manager* of the department which controls these project(s);

Associate **EMPLOYEE** with **DEPARTMENT** to get the manager, and *then* **DEPARTMENT** with **PROJECT** to get the controlled projects by this department.

```
( SELECT      DISTINCT Pnumber
  FROM        PROJECT, DEPARTMENT, EMPLOYEE
 WHERE        Dnum=Dnumber AND Mgr_ssn=Ssn
              AND Lname='Smith' )
```



IN-CLASS ACTIVITY [A2]

Step 3: UNION over the two sets of project numbers:

```
( SELECT      DISTINCT Pnumber
  FROM        PROJECT, DEPARTMENT, EMPLOYEE
  WHERE       Dnum=Dnumber AND Mgr_ssn=Ssn
              AND Lname='Smith' )

UNION

( SELECT      DISTINCT Pnumber
  FROM        PROJECT, WORKS_ON, EMPLOYEE
  WHERE       Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```

THREE-VALUED LOGIC

SQL is a three-valued logic: **TRUE** (1), **FALSE** (0) and **UNKNOWN** (0.5)

Recall: Each NULL value is *different* from any other NULL value!

Principle: Any value compared with NULL evaluates to **UNKNOWN**

Example:

WHERE Address = NULL ...evaluates to **UNKNOWN**;

WHERE Address <> NULL ...evaluates to **UNKNOWN**;

WHERE NULL = NULL ...evaluates to **UNKNOWN**

Always adopt: IS NULL or IS NOT NULL

AND	TRUE	FALSE	UNKNOWN
TRUE	TRUE	FALSE	UNKNOWN
FALSE	FALSE	FALSE	FALSE
UNKNOWN	UNKNOWN	FALSE	UNKNOWN

min

OR	TRUE	FALSE	UNKNOWN
TRUE	TRUE	TRUE	TRUE
FALSE	TRUE	FALSE	UNKNOWN
UNKNOWN	TRUE	UNKNOWN	UNKNOWN

max

NOT	
TRUE	FALSE
FALSE	TRUE
UNKNOWN	UNKNOWN

$1-x$



COMPARISON INVOLVING NULL

Query 6: Retrieve the first and last names of *all* employees who do not have supervisors.

```
SELECT    Fname, Lname
FROM      EMPLOYEE
WHERE     Super_ssn IS NULL;
```

```
SELECT Fname, Lname
FROM   EMPLOYEE
WHERE  Super_ssn = NULL
```

} it produces *no* tuples!
Hence, *wrong* reasoning!
Why?



IN-CLASS ACTIVITY [A3]

```
CREATE TABLE Recruitment (  
  ID INT NOT NULL,  
  Name VARCHAR(10),  
  RecruiterID INT,  
  PRIMARY KEY(ID),  
  FOREIGN KEY(RecruiterID)  
  REFERENCES Recruitment(ID)  
);
```

ID	Name	RecruiterID
1	Chris	NULL
2	Philip	3
3	Stella	NULL
4	John	2
5	Teresa	3
6	Iona	5

Task 1: List *all* the names of the employees.

Task 2: List the names of the employees who are *not* recruited by employee with ID = 2.

Which is the problem here?

IN-CLASS ACTIVITY [A3]



{Chris, Philip, Stella, John, Teresa, Iona}

```
SELECT Name, RecruiterID
FROM Recruitment
WHERE RecruiterID <> 2
{Philip, Teresa, Iona}
```

We *imply*...{Chris, Stella, John} are recruited by Philip.

```
SELECT Name, RecruiterID
FROM Recruitment
WHERE RecruiterID IS NULL
      OR RecruiterID <> 2
{Chris, Philip, Stella, Teresa, Iona}
```

ID	Name	RecruiterID
1	Chris	NULL
2	Philip	3
3	Stella	NULL
4	John	2
5	Teresa	3
6	Iona	5

NESTED (INNER) QUERY



Nested query is a query *within* another (*outer*) query;

- SELECT-FROM-WHERE block *within* another outer WHERE clause.
- Nested query's **output** is **input** to outer's WHERE via: **IN**, **ALL**, **EXISTS**
- **Nested Uncorrelated Query**: *first* execute the nested query, and *then* execute the outer query using inner's output.
- **Correlated Query**: for *each* tuple of the outer query, we execute the nested query.



NESTED UNCORRELATED QUERY: OPERATOR **IN**



IN: checks whether a value belongs to the inner's output set (or *multiset*), i.e.,

Query 7: Show the SSN of those employees who work in the projects with number: either 1, or 2, or 3.

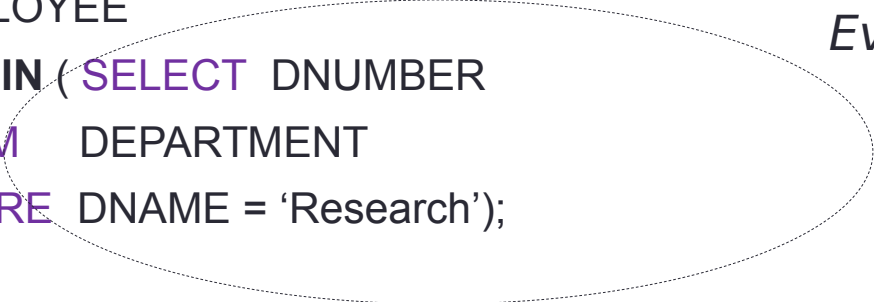
```
SELECT Essn
FROM   WORKS_ON
WHERE  PNO IN (1, 2, 3);
```

- if PNO = 1 then PNO **IN** (1, 2, 3) evaluates to **TRUE**
- if PNO = 4 then PNO **IN** (1, 2, 3) evaluates to **FALSE**

NESTED UNCORRELATED QUERY: OPERATOR **IN**

Query 8: Show the names of those employees who work in the department 'Research'.

```
SELECT FNAME  
FROM EMPLOYEE  
WHERE DNO IN ( SELECT DNUMBER  
                FROM DEPARTMENT  
                WHERE DNAME = 'Research');
```



Evaluates to **5**

```
SELECT FNAME  
FROM EMPLOYEE  
WHERE DNO IN ( 5 );
```

NESTED UNCORRELATED QUERY: OPERATOR **ALL**



ALL: compares a value with *all* the values from the inner's output set.

Query 9: Show the last and first names of those employees whose salary is *greater* than the salaries of *all* employees in Department 5.

```
SELECT      Lname, Fname
FROM        EMPLOYEE
WHERE       Salary > ALL ( SELECT      Salary
                           FROM        EMPLOYEE
                           WHERE       Dno=5 );
```

} First, find *all* salaries from employees in Department 5;

NESTED CORRELATED QUERY

For **each** *tuple* of the *outer* query, we execute the inner query!

Relation as a variable: **global scope** (*outer*) and **local scope** (*inner*).

Query 10: Retrieve the name of *each* employee who has a dependent with the same first name and is the same gender as that employee.

SELECT
FROM
WHERE

E.Fname, E.Lname

EMPLOYEE AS E

E.Ssn IN (SELECT

FROM
WHERE

Essn

DEPENDENT AS D

E.Fname=D.Dependent_name
AND E.Sex=D.Sex);

For each *outer*
employee E, retrieve
the dependents D
and check!



NESTED CORRELATED QUERY

Lemma 1: Correlated queries using **IN** can be collapsed into one *single* block.

Query 11: Retrieve the name of *each* employee who has a dependent with the same first name and is the same gender as that employee.

```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE AS E, DEPENDENT AS D
WHERE  E.Ssn=D.Essn
       AND E.Sex=D.Sex
       AND E.Fname=D.Dependent_name;
```

NESTED CORRELATED QUERY: EXISTS



EXISTS: checks *whether* the inner's output is *empty* set or *not*, and returns FALSE or TRUE, respectively, e.g., or

Opposite: NOT EXISTS

```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE AS E
WHERE  EXISTS
      (SELECT * FROM DEPARTMENT AS D WHERE E.DNO = D.DNUMBER)
```

- Checks if a *given* employee is working at *some* department.
- Reason about **E.DNO** being NULL.

NESTED CORRELATED QUERY: EXISTS



```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE AS E
WHERE  EXISTS
      (SELECT *
       FROM   DEPARTMENT AS D
       WHERE  E.DNO = D.DNUMBER AND D.DNAME = 'Research' )
```

➤ Describe this...

IN-CLASS QUIZ



```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE AS E
WHERE  EXISTS
  (SELECT * FROM DEPENDENT AS P WHERE E.SSN = P.Essn)
  AND EXISTS
  (SELECT * FROM DEPARTMENT AS D WHERE E.SSN = D.Mgr_SSN)
```

Checks if a *given* employee:

- has at least a dependent **and**
- manages a department, i.e., there *exists* a department, which is managed by that employee.

IN-CLASS ACTIVITY [A4]



STUDENT (Name, StudentID, Class)

COURSE (Name, CourseID, Credits, School)

GRADES (StudentID, CourseID, Grade)

*/*Grade: {'A', 'B', 'C', 'D', 'E'}*/*

Task: Retrieve the names of *all* students who have a Grade of 'A' in *all* of their courses.

IN-CLASS ACTIVITY [A4]



STUDENT (Name, StudentID, Class)

COURSE (Name, CourseID, Credits, School)

GRADES (StudentID, CourseID, Grade)

/*Grade: {'A', 'B', 'C', 'D', 'E'}*/

```
SELECT  S.Name
FROM    STUDENT S
WHERE   NOT EXISTS
        (SELECT * FROM GRADES G
         WHERE G.StudentID = S.StudentID
              AND NOT (G.Grade = 'A')
        )
```

There does
not exist a
grade which
is *not* 'A', i.e.,
all grades are
'A'



IN-CLASS QUIZ [WHAT WE GET]

```
SELECT *  
FROM EMPLOYEE  
WHERE EXISTS (SELECT NULL FROM EMPLOYEE);
```

```
SELECT *  
FROM EMPLOYEE  
WHERE EXISTS (SELECT * FROM EMPLOYEE  
              WHERE NULL IS NULL);
```

Fact: SQL, in evaluating EXISTS, *simply counts* rows and *ignores* the value(s) in the subquery—even if they are NULL!



IN-CLASS QUIZ [WHAT WE GET]

```
SELECT * FROM EMPLOYEE WHERE  
    EXISTS (SELECT 0  
        FROM EMPLOYEE  
            WHERE 1 IS NULL OR NULL);
```

```
SELECT * FROM EMPLOYEE WHERE  
    NOT EXISTS (SELECT 0  
        FROM EMPLOYEE  
            WHERE 1 IS NOT NULL OR NULL);
```