

# Pandas: Data Structures, Reading and Writing

# Basic introduction

- pandas (Python Data Analysis Library) is a library specialized for data analysis and used a series of I/O API functions to read and write data as dataframe objects.
- Python IDE (Integrated Development Environment): Jupyter Notebook or Spyder.
- Easiest way to get Pandas set up is to install it through a package like the Anaconda distribution.
- Primary data structures on which all transactions are centred (generally made during the analysis):
  - **Series**: Object of the library designed to represent one-dimensional data structure.
  - **Dataframes**: More complex data structure designed to contain cases with several dimensions.

# The Series

```
import pandas as pd
import numpy as np
```

```
s = pd.Series([12, -4, 7, 9])
```

← Declaring a series

```
0    12
1    -4
2     7
3     9
dtype: int64
```

```
s[2]
```

← Selecting one internal element

```
7
```

```
s[0:3]
```

← Selecting multiple elements

```
0    12
1    -4
2     7
dtype: int64
```

```
s[1] = 0
```

← Assigning a value to an item using its index.

```
0    12
1     0
2     7
3     9
dtype: int64
```

```
s = pd.Series([12, -4, 7, 9],
              index=['a', 'b', 'c', 'd'])
```

← Declaring a series, assigning an index

```
s
a     12
b     -4
c      7
d      9
dtype: int64
```

```
s['b'] = 100
```

← Assigning a value to an item using its label

```
s
a     12
b    100
c      7
d      9
dtype: int64
```

```
s[s > 8]
```

← Filtering values

```
a     12
b    100
d      9
dtype: int64
```

```
s / 2
```

← Operators (+, -, \*, and /) and mathematical function that are applicable to NumPy array can be extended to Series

```
a     6.0
b    -2.0
c     3.5
d     4.5
```

# The Series – Evaluating Values

```
serd = pd.Series([1,0,2,1,2,3],  
                  index=['white','white','blue',  
                        'green','green','yellow'])  
serd
```

```
white    1  
white    0  
blue     2  
green    1  
green    2  
yellow   3  
dtype: int64
```

```
serd.unique()
```

```
array([1, 0, 2, 3], dtype=int64)
```

```
serd.value_counts()
```

```
2    2  
1    2  
3    1  
0    1  
dtype: int64
```

```
serd.isin([0,3])
```

```
white    False  
white     True  
blue     False  
green    False  
green    False  
yellow    True  
dtype: bool
```

- The `unique()` function will tell us all the values contained in a series, excluding duplicates
- The `value_counts()` will return the unique values but also calculate the occurrences within a series.
- The `isin()` function tells us if the values are contained in the data structure. Boolean values returned can be very useful when filtering data in a series or in a column of a dataframe.

## The Series – NaN Values

- **NaN** (Not a Number) is used in pandas data structures to indicate the presence of an empty field or a non-numeric element. To define a missing value, we enter `np.NaN`.
- The `isnull()` and `notnull()` functions are useful to identify the indexes without a value.
- The `isnull()` returns `True` at NaN values in the series.
- The `notnull()` returns `True` if they are not NaN.
- These functions are often put inside filters to make a condition

```
s2 = pd.Series([5, -3, np.NaN, 20])  
s2  
  
0      5.0  
1     -3.0  
2      NaN  
3     20.0  
dtype: float64
```

```
s2.isnull()
```

```
0      False  
1      False  
2         True  
3      False  
dtype: bool
```

```
s2.notnull()
```

```
0      True  
1      True  
2     False  
3      True  
dtype: bool
```

```
s2[s2.notnull()]
```

```
0      5.0  
1     -3.0  
3     20.0  
dtype: float64
```

```
s2[s2.isnull()]
```

```
2      NaN  
dtype: float64
```

## The Series

We can create a series from a previously defined dictionary. The array of the index is filled with the keys while the data are filled with their values.

```
mydict = {'red':250,'blue': 560,  
          'green':700,'white':1456}  
mydict  
{'red': 250, 'blue': 560, 'green': 700, 'white': 1456}
```

```
myseries = pd.Series(mydict)  
myseries  
red      250  
blue     560  
green    700  
white   1456  
dtype: int64
```

We can also define the array indexes separately. As seen, if there is a mismatch, pandas will add the NaN value.

```
colours = ['red','blue','green','white', 'purple']  
myseries = pd.Series(mydict, index = colours)  
myseries
```

```
red      250.0  
blue     560.0  
green    700.0  
white   1456.0  
purple      NaN  
dtype: float64
```

We can perform operations between two series. Series can align data addressed differently between them by identifying their corresponding labels.

```
mydict2 = {'red':900,'black':800,'white':500}  
myseries2 = pd.Series(mydict2)  
myseries + mysereis2
```

```
black      NaN  
blue       NaN  
green      NaN  
purple     NaN  
red      1150.0  
white    1956.0  
dtype: float64
```

# The DataFrame

- Tabular structure very similar to a spreadsheet.
- Designed to extend series to multiple dimensions.
- Consists of an ordered collection of columns, each of which can contain a value of a different type (numeric, string, Boolean, etc.)
- Unlike series which have an index array containing labels associated with each element, the dataframe has two index arrays.
- It can be understood as a dictionary of series where the keys are the column names and the values are the series that will form the columns of the dataframe

```
import numpy as np
import pandas as pd
```

```
data = {'color': ['white', 'red', 'black', 'green', 'purple'],
        'items': ['ball', 'pen', 'pencil', 'paper', 'eraser'],
        'price': [2.5, 1.5, 0.5, 0.6, 0.15]}
```

```
frame = pd.DataFrame(data)
frame
```

Define a  
dataframe

	color	items	price
0	white	ball	2.50
1	red	pen	1.50
2	black	pencil	0.50
3	green	paper	0.60
4	purple	eraser	0.15

```
frame2 = pd.DataFrame(data, columns=['items', 'price'])
frame2
```

	items	price
0	ball	2.50
1	pen	1.50
2	pencil	0.50
3	paper	0.60
4	eraser	0.15

We can select the data we want to display. Use column option to specify the sequence of columns. dataframe

# The DataFrame

- A common data structure used in Python is a `nested dict`. When it is passed directly as an argument to the `DataFrame()` constructor, pandas will treat external keys as column names and internal keys as labels for indexes.
- Fields with no match are assigned the `NaN` value.
- **Transposition**: columns become row and rows become columns. It is achieved by adding the `T` attribute to its operation

```
nesteddict = {'red':{2012:22, 2014:45},  
             'green':{2008: 23, 2012:22,2014:17},  
             'blue':{2008: 18,2012:28,2014: 19}}  
frame3 = pd.DataFrame(nesteddict)  
frame3
```

	red	green	blue
2008	NaN	23	18
2012	22.0	22	28
2014	45.0	17	19

frame3.T

	2008	2012	2014
red	NaN	22.0	45.0
green	23.0	22.0	17.0
blue	18.0	28.0	19.0



## Reading Data in CSV or Text files

- Most common operation for data analysis is to read the data contained in a .CSV file or even a text file.
- To achieve the, we must import the following libraries `numpy` and `pandas` in our Jupyter Notebook
- The `read_csv()` function will read the content of the `.csv` file and convert it to a dataframe object.

```
In [6]: import numpy as np
import pandas as pd
csvframe = pd.read_csv('Documents/texting.csv')
csvframe
```

Out[6]:


	white	red	blue	green	animal
0	1	5	2	3	car
1	2	7	8	5	dog
2	3	3	6	7	horse
3	2	2	8	3	duck
4	4	4	2	1	mouse

```
In [13]: csvframe1 = pd.read_csv('Documents/texting.txt')
csvframe1
```

Out[13]:

	white	red	blue	green	animal
0	1	5	2	3	cat
1	2	7	8	5	dog
2	3	3	6	7	horse
3	2	2	8	3	duck
4	4	4	2	1	mouse

Original texting.csv and  
texting.txt files as seen  
in the spreadsheet and  
notepad



white	red	blue	green	animal
1	5	2	3	car
2	7	8	5	dog
3	3	6	7	horse
2	2	8	3	duck
4	4	2	1	mouse

```
white,red,blue,green,animal
1,5,2,3,cat
2,7,8,5,dog
3,3,6,7,horse
2,2,8,3,duck
4,4,2,1,mouse
```

## Reading Data in CSV or Text files

- When using `read_table()` function to read a csv or txt file, specify the delimiter otherwise, the data will not be in a tabulated format.

Output of `read_table()` function without specifying the delimiter

In [15]: `pd.read_table('Documents/texting.csv')`

Out[15]:

	white,red,blue,green,animal
0	1,5,2,3,car
1	2,7,8,5,dog
2	3,3,6,7,horse
3	2,2,8,3,duck
4	4,4,2,1,mouse

Output of `read_table()` function with specified delimiter

In [16]: `pd.read_table('Documents/texting.csv', sep=',')`

Out[16]:

	white	red	blue	green	animal
0	1	5	2	3	car
1	2	7	8	5	dog
2	3	3	6	7	horse
3	2	2	8	3	duck
4	4	4	2	1	mouse

# Reading Data in CSV or Text files

- `pd.read_csv('Documents/texting.csv', header=None)`: This will tell pandas to assign the default name to the columns.
- You can specify the names directly by assigning a list of labels to the `names` options `pd.read_csv('Documents/texting.csv', names=['white', 'red', 'blue', 'green', 'animal'])`
- Create a dataframe with a hierarchical structure by extending the functionality of the `read_csv()` function by adding the `index_col` option.

Original Hierarchical data

```
In [22]: pd.read_csv('Documents/texting1.txt', index_col=['color', 'status'])
```

Out[22]:

		item1	item2	item3
color	status			
black	up	3	4	6
	down	2	6	7
white	up	5	5	5
	down	3	3	2
	left	1	2	1
red	up	2	2	2
	down	1	1	4

texting1 - Notepad

File Edit Format View Help

```
color,status,item1,item2,item3
black,up,3,4,6
black,down,2,6,7
white,up,5,5,5
white,down,3,3,2
white,left,1,2,1
red,up,2,2,2
red,down,1,1,4
```

# Using RegExp to Parse TXT files

- Sometimes the files on which to parse the data do not show separators such as comma or a semicolon.
- Regular expressions can be used as criteria for value separation.

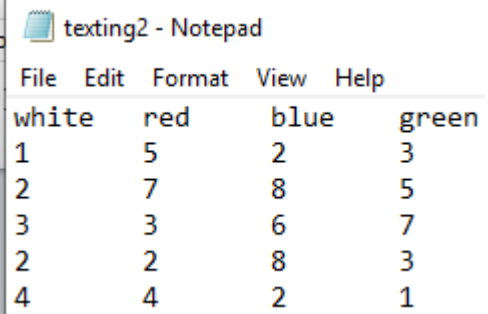
.	Single character, except newline
\d	Digit
\D	Non-digit character
\s	Whitespace character
\S	Non-whitespace character
\n	New line character
\t	Tab character
\uxxxx	Unicode character specified by the hexadecimal number xxxx

# Using RegExp to Parse TXT files - Examples

```
In [23]: pd.read_table('Documents/texting2.txt', sep='\s+', engine='python')
```

Out[23]:

	white	red	blue	green
0	1	5	2	3
1	2	7	8	5
2	3	3	6	7
3	2	2	8	3
4	4	4	2	1



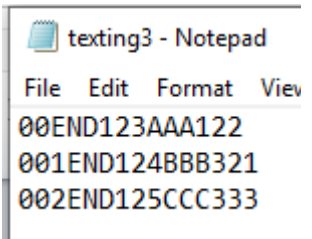
texting2 - Notepad

File	Edit	Format	View	Help
white	red	blue	green	
1	5	2	3	
2	7	8	5	
3	3	6	7	
2	2	8	3	
4	4	2	1	

```
In [28]: pd.read_table('Documents/texting3.txt', sep='\D+', header=None, engine='python')
```

Out[28]:

	0	1	2
0	0	123	122
1	1	124	321
2	2	125	333



texting3 - Notepad

File	Edit	Format	View
00END123AAA122			
001END124BBB321			
002END125CCC333			

This example extract the numeric part from a texting3.txt file. Header option = None because the column heading is not in the texting3.txt file.

# Using RegExp to Parse TXT files - Examples

With the `skiprows` option, you can exclude the lines you want.

```
pd.read_table('Documents/texting2.txt', sep='\D+', header=None, engine='python', skiprows=[0,1,2])
```

	0	1	2	3
0	3	3	6	7
1	2	2	8	3
2	4	4	2	1

texting2.txt file  
dataframe after  
skipping the first  
three rows

	0	1	2	3
0	white	red	blue	green
1	1	5	2	3
2	2	7	8	5
3	3	3	6	7
4	2	2	8	3
5	4	4	2	1

Original  
texting2.txt  
file  
dataframe

# Reading TXT Files into Parts

- To read only a portion of the file, you can specify the numbers of lines on which to parse. You can use the `nrows` and `skiprows` options.

```
pd.read_csv('Documents/texting.csv', skiprows=[2], nrows=3, header=None)
```

	0	1	2	3	4
0	white	red	blue	green	animal
1	1	5	2	3	car
2	3	3	6	7	horse

Dataframe after using  
`nrows` and `skiprows`

	0	1	2	3	4
0	white	red	blue	green	animal
1	1	5	2	3	car
2	2	7	8	5	dog
3	3	3	6	7	horse
4	2	2	8	3	duck
5	4	4	2	1	mouse

Original  
dataframe

# Writing Data in CSV

- `to_csv()` : Function used to write the data contained in a dataframe to a csv file.

```
frame = pd.DataFrame(np.arange(16).reshape((4,4)),  
                      index = ['red', 'blue', 'yellow', 'white'],  
                      columns = ['ball', 'pen', 'pencil', 'paper'])
```

```
frame.to_csv('writin.txt')
```

```
pd.read_csv('writin.txt')
```

Unnamed: 0	ball	pen	pencil	paper	
0	red	0	1	2	3
1	blue	4	5	6	7
2	yellow	8	9	10	11
3	white	12	13	14	15

The `index` and `header` to `False` options are used to remove the default indexes and columns that are marked on the file by default.

```
frame.to_csv('writin.txt', index=False, header=False)
```

	0	1	2	3
0	4	5	6	7
1	8	9	10	11
2	12	13	14	15



# Reading and Writing HTML & Excel Files

- `to_html()` function will convert a dataframe into an HTML table.
- `read_html()` function returns a list of dataframes even if there is only one table.
- `to_excel()` function to convert a dataframe into a spreadsheet on Excel.
- `read_excel()` read the data contained in the excel file and convert it into a dataframe

```
frame = pd.DataFrame(np.arange(4).reshape(2,2))
```

```
print(frame.to_html())
```

```
<table border="1" class="dataframe">
  <thead>
    <tr style="text-align: right;">
      <th></th>
      <th>0</th>
      <th>1</th>
    </tr>
  </thead>
  <tbody>
    <tr>
      <th>0</th>
      <td>0</td>
      <td>1</td>
    </tr>
    <tr>
      <th>1</th>
      <td>2</td>
      <td>3</td>
    </tr>
  </tbody>
</table>
```

# Interacting with Databases

- `pandas.io.sql` module provides a unified interface independent of the DB called `sqlalchemy`. This interface simplifies the connection mode, regardless of the commands will be always be the same.
- The `create_engine()` function is used to make a connection.
- Example of code for connecting different databases.

```
For PostgreSQL
engine = create_engine('postgresql://mireilla:text@localhost:5432/mydatabase')

For MySQL
engine = create_engine('mysql+mysqldb://mireilla:text@localhost/foo')

For Oracle
engine = create_engine('oracle://mireilla:text@127.0.0.1:1521/sidname')

For MSSQL
engine = create_engine('mssql+pyodbc://mydsn')

For SQLite
engine = create_engine('sqlite:///foo.db')
```

# Interacting with Databases – SQLite3

```
from sqlalchemy import create_engine
```

```
frame = pd.DataFrame(np.arange(20).reshape(4,5),  
                     columns=['white', 'red', 'blue', 'black', 'green'])
```

frame

	white	red	blue	black	green
0	0	1	2	3	4
1	5	6	7	8	9
2	10	11	12	13	14
3	15	16	17	18	19

```
engine = create_engine('sqlite:///foo.db')
```

```
frame.to_sql('colors',engine)
```

```
pd.read_sql('colors',engine)
```

	index	white	red	blue	black	green
0	0	0	1	2	3	4
1	1	5	6	7	8	9
2	2	10	11	12	13	14
3	3	15	16	17	18	19

Create a dataframe that will be used to create a new table on the SQLite3 database

Implement the connection to the SQLite3 database

Convert the dataframe

Read the database with the read\_sql() function with the name of the table and the engine