

idss_lecture_01_introduction_demo

September 20, 2022

```
[1]: %%javascript
    $('#menubar').toggle();
```

<IPython.core.display.Javascript object>

1 Introduction to Data Science and Systems 2022-2023

1.1 Lecture Week 1: Introduction - *a basic data science example*

University of Glasgow v20222023a - Example adapted from A. Joseph and the *DS100 textbook*.

1.2 Load the basic packages...

```
[2]: !pip install plotly
```

```
Requirement already satisfied: plotly in
/home/nicolas/anaconda3/lib/python3.9/site-packages (5.6.0)
Requirement already satisfied: tenacity>=6.2.0 in
/home/nicolas/anaconda3/lib/python3.9/site-packages (from plotly) (8.0.1)
Requirement already satisfied: six in
/home/nicolas/anaconda3/lib/python3.9/site-packages (from plotly) (1.16.0)
```

```
[3]: # You do not need to understand all of this in detail yet
    # - this notebook is just for inspiration and motivation

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns # (this is a fancy plotting library - we will mostly be
    ↪ using matplotlib in this course)

## Plotly plotting support (this is a fancy plotting library - we will mostly be
    ↪ using matplotlib in this course)
import plotly.offline as py
py.init_notebook_mode()
import plotly.graph_objs as go
```

```
import plotly.figure_factory as ff
import plotly.express as px
```

1.3 1) The questions

- How many students in IDSS this semester ?
 - What is the distribution of the degree programmes in IDSS this year?
 - Key questions: What is the gender distribution in IDSS this semester?
-

1.4 2) Data acquisition and availability

- You provided some of the data when registering via MyCampus. The data was thus collected for general purpose administrating and not necessarily to answer your specific question.
-

1.5 3) Storage and access

- MyCampus (Oracle)
- SoCS LTC system (SQL)
- Moodle (activity)

Some observations: - The data based contain sensitive information - we need to make sure the data is safely and can only be accessed by approved users (and inline with GDPR)

2 4) Data exploration, preparation and curation

2.1 4.1) Loading and inspecting the data using Pandas

We need a data structure to load the data into for visualisation, querying and exploration.

Here we use Pandas because they provide built-in functionality to easily explore the (see Moodle for suggested study material)

You can read more about panda here <https://pandas.pydata.org/>

```
[4]: data = pd.read_csv("roster_mod_20202021c.csv")
len(data)
```

```
[4]: 374
```

```
[5]: data.head(20)
```

```
[5]:   Level ProgrammeID  Firstname  Advisor \
0      UG          NaN    Andrew  Rogers,S
1      PGT    I261-5200    Mitko   Rogers,S
```

2	PGT	I261-5200	Euan	Rogers,S
3	PGT	G511-5200	John	Rogers,S
4	PGT	I261-5200	Lucas	Rogers,S
5	PGT	G577-5200	Grant	ChiefAdviser-Science,0
6	PGT	G577-5200	Kleanthis	Rogers,S
7	PGT	I261-5200	Charles	Rogers,S
8	PGT	I261-5200	Xiao	Rogers,S
9	PGT	I261-5200	Pengjun	Rogers,S
10	PGT	I261-5200	Qiling	Rogers,S
11	PGT	I261-5200	Yuli	Rogers,S
12	PGT	G511-5200	Jihua	Rogers,S
13	PGT	I261-5200	Mengting	ChiefAdviser-Science,0
14	PGT	I261-5200	Yuchen	Rogers,S
15	PGT	G577-5200	Zhehao	Rogers,S
16	PGT	I261-5200	Zitong	Rogers,S
17	PGT	I261-5200	Guanxuan	ChiefAdviser-Science,0
18	PGT	I261-5200	YINGHONG	Rogers,S
19	PGT	G511-5200	Mohammad	Rogers,S

	Programme	Gender
0	Computer Science, BSc	M
1	MSc in Data Science	NaN
2	MSc in Data Science	NaN
3	Computing Science,MSc	NaN
4	MSc in Data Science	NaN
5	Information Security,MSc	NaN
6	Information Security,MSc	NaN
7	MSc in Data Science	NaN
8	MSc in Data Science	NaN
9	MSc in Data Science	NaN
10	MSc in Data Science	NaN
11	MSc in Data Science	NaN
12	Computing Science,MSc	NaN
13	MSc in Data Science	NaN
14	MSc in Data Science	NaN
15	Information Security,MSc	NaN
16	MSc in Data Science	NaN
17	MSc in Data Science	NaN
18	MSc in Data Science	NaN
19	Computing Science,MSc	NaN

```
[6]: data['Firstname'].unique()
len(data['Firstname'].unique())
```

[6]: 349

2.2 4.2) Curation and cleaning

```
[7]: data['Firstname'] = data['Firstname'].str.lower()
print("Number of Students:", len(data))
data.head(20)
```

Number of Students: 374

```
[7]:
```

	Level	ProgrammeID	Firstname	Advisor \
0	UG	NaN	andrew	Rogers,S
1	PGT	I261-5200	mitko	Rogers,S
2	PGT	I261-5200	euan	Rogers,S
3	PGT	G511-5200	john	Rogers,S
4	PGT	I261-5200	lucas	Rogers,S
5	PGT	G577-5200	grant	ChiefAdviser-Science,0
6	PGT	G577-5200	kleanthis	Rogers,S
7	PGT	I261-5200	charles	Rogers,S
8	PGT	I261-5200	xiao	Rogers,S
9	PGT	I261-5200	pengjun	Rogers,S
10	PGT	I261-5200	qiling	Rogers,S
11	PGT	I261-5200	yuli	Rogers,S
12	PGT	G511-5200	jihua	Rogers,S
13	PGT	I261-5200	mengting	ChiefAdviser-Science,0
14	PGT	I261-5200	yuchen	Rogers,S
15	PGT	G577-5200	zhehao	Rogers,S
16	PGT	I261-5200	zitong	Rogers,S
17	PGT	I261-5200	guanxuan	ChiefAdviser-Science,0
18	PGT	I261-5200	yinghong	Rogers,S
19	PGT	G511-5200	mohammad	Rogers,S

	Programme	Gender
0	Computer Science, BSc	M
1	MSc in Data Science	NaN
2	MSc in Data Science	NaN
3	Computing Science,MSc	NaN
4	MSc in Data Science	NaN
5	Information Security,MSc	NaN
6	Information Security,MSc	NaN
7	MSc in Data Science	NaN
8	MSc in Data Science	NaN
9	MSc in Data Science	NaN
10	MSc in Data Science	NaN
11	MSc in Data Science	NaN
12	Computing Science,MSc	NaN
13	MSc in Data Science	NaN
14	MSc in Data Science	NaN
15	Information Security,MSc	NaN

```

16      MSc in Data Science      NaN
17      MSc in Data Science      NaN
18      MSc in Data Science      NaN
19      Computing Science,MSc     NaN

```

```

[8]: data['Firstname'].unique()
len(data['Firstname'].unique())

```

```

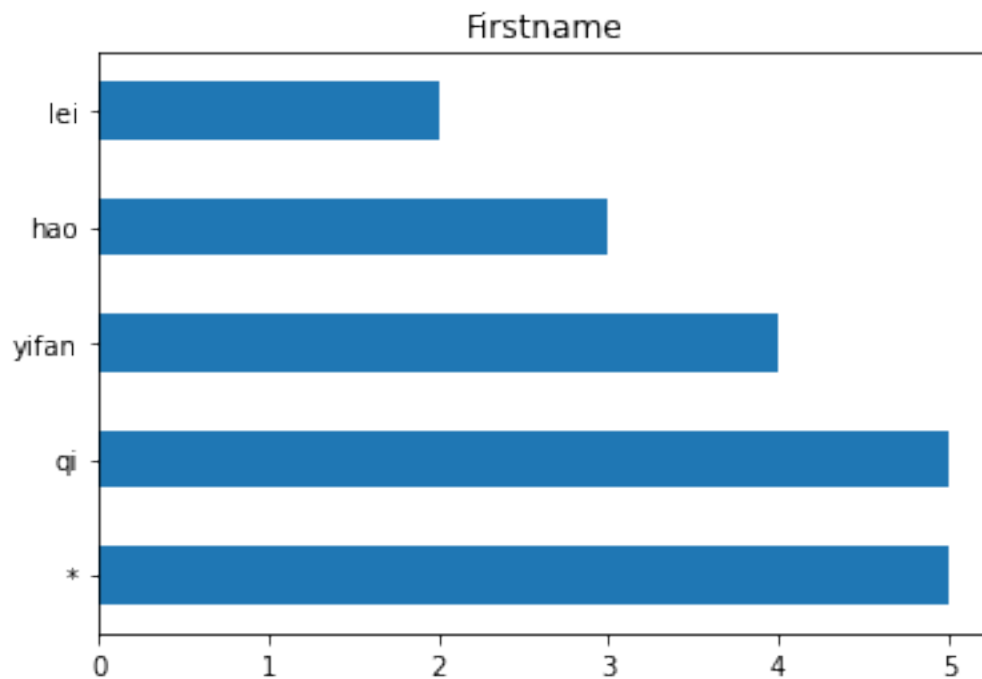
[8]: 343

```

```

[9]: (
    data["Firstname"]
      .str.lower()
      .value_counts().sort_values(ascending=False)
      .head(5).plot(kind='barh', title = "Firstname")
);

```



```

[10]: data = data.drop(data[data.Programme == "Computer Science, BSc"].index)
len(data)

```

```

[10]: 372

```

2.3 4.3) Basic visualization and summarization

```
[11]: data.describe()
```

```
[11]:
```

	Level	ProgrammeID	Firstname	Advisor	Programme	Gender
count	372	372	372	372	370	6
unique	1	3	341	12	3	2
top	PGT	I261-5200	*	Rogers,S	MSc in Data Science	F
freq	372	190	5	56	188	4

3 5) Data modelling and analysis

3.1 Q: What is the distribution of the lengths of students' names in this class?


```
[47]: sns.distplot(data['Firstname'].str.len(), rug=True, axlabel="Number of ↵  
      Characters");
```

```
/home/nicolas/anaconda3/lib/python3.9/site-  
packages/seaborn/distributions.py:2619: FutureWarning:
```

```
`distplot` is a deprecated function and will be removed in a future version.  
Please adapt your code to use either `displot` (a figure-level function with  
similar flexibility) or `histplot` (an axes-level function for histograms).
```

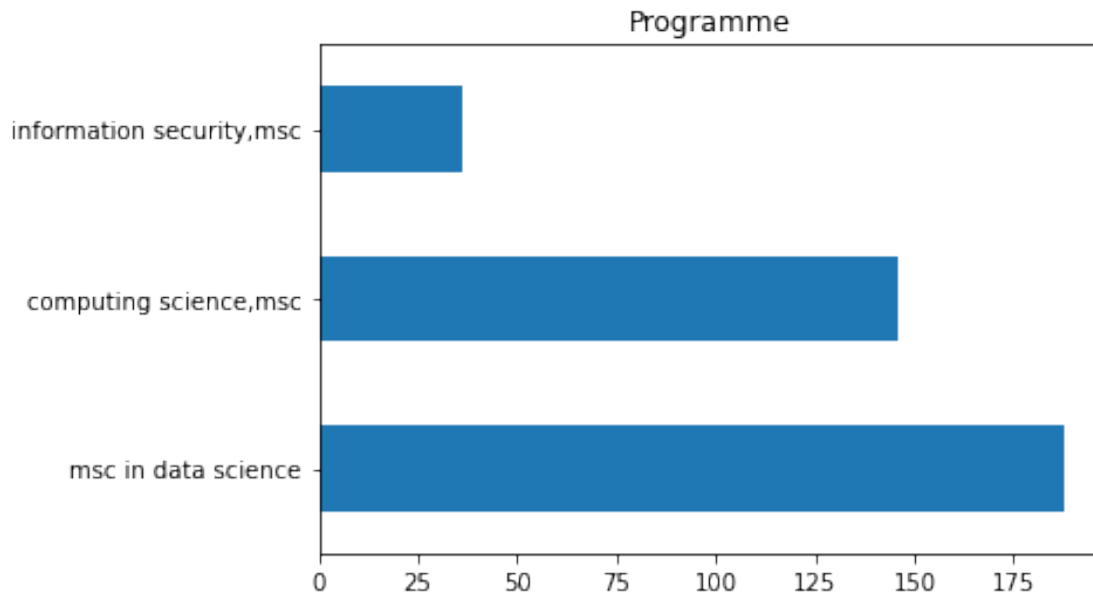
```
/home/nicolas/anaconda3/lib/python3.9/site-  
packages/seaborn/distributions.py:2103: FutureWarning:
```

```
The `axis` variable is no longer used and will be removed. Instead, assign  
variables directly to `x` or `y`.
```




3.2 Q: What are the programme degrees of students in the class?

```
[13]: (  
    data["Programme"]  
      .str.lower()  
      .value_counts().sort_values(ascending=False)  
      .head(20).plot(kind='barh', title = "Programme")  
);
```



```
[14]: # fancy plotting using plotly (we won't actually use it the course; is only for
      ↳ inspiration)
      px.bar(data['Programme'].value_counts().to_frame().reset_index().head(20),
             x = 'Programme',
             y = 'index',
             orientation = 'h')
```



3.3 Q: What is the gender distribution of the class?

Can we answer this questions with the raw data?

```
[15]: print(data.columns)
```



```
Index(['Level', 'ProgrammeID', 'Firstname', 'Advisor', 'Programme', 'Gender'],
      dtype='object')
```

So the answer is maybe as we do have a column called Gender, but remember all the NaNs...

```
[16]: data.head(20)
```

```
[16]:
```

	Level	ProgrammeID	Firstname	Advisor	\
1	PGT	I261-5200	mitko	Rogers,S	
2	PGT	I261-5200	euan	Rogers,S	
3	PGT	G511-5200	john	Rogers,S	
4	PGT	I261-5200	lucas	Rogers,S	
5	PGT	G577-5200	grant	ChiefAdviser-Science,0	
6	PGT	G577-5200	kleanthis	Rogers,S	
7	PGT	I261-5200	charles	Rogers,S	
8	PGT	I261-5200	xiao	Rogers,S	
9	PGT	I261-5200	pengjun	Rogers,S	
10	PGT	I261-5200	qiling	Rogers,S	
11	PGT	I261-5200	yuli	Rogers,S	
12	PGT	G511-5200	jihua	Rogers,S	
13	PGT	I261-5200	mengting	ChiefAdviser-Science,0	
14	PGT	I261-5200	yuchen	Rogers,S	
15	PGT	G577-5200	zhehao	Rogers,S	
16	PGT	I261-5200	zitong	Rogers,S	
17	PGT	I261-5200	guanxuan	ChiefAdviser-Science,0	
18	PGT	I261-5200	yinghong	Rogers,S	
19	PGT	G511-5200	mohammad	Rogers,S	
20	PGT	I261-5200	yuzhou	Rogers,S	

	Programme	Gender
1	MSc in Data Science	NaN
2	MSc in Data Science	NaN
3	Computing Science,MSc	NaN
4	MSc in Data Science	NaN
5	Information Security,MSc	NaN
6	Information Security,MSc	NaN
7	MSc in Data Science	NaN
8	MSc in Data Science	NaN
9	MSc in Data Science	NaN
10	MSc in Data Science	NaN
11	MSc in Data Science	NaN
12	Computing Science,MSc	NaN
13	MSc in Data Science	NaN
14	MSc in Data Science	NaN
15	Information Security,MSc	NaN
16	MSc in Data Science	NaN
17	MSc in Data Science	NaN
18	MSc in Data Science	NaN

19	Computing Science, MSc	NaN
20	MSc in Data Science	NaN

Ideas:

- What do we mean by gender?
- Can we use the name to estimate gender?
- How would we build model of gender given the name?
- Where can we get data for such a model?

3.3.1 Build a model based on an aux dataset

- Easily accessible data is from the US SSA (<https://www.ssa.gov/oact/babynames/>)? *Is it going to cause problem that we are using US data in the UK?*

```
[17]: import urllib.request
import os.path

# Download data from the web directly
data_url = "https://www.ssa.gov/oact/babynames/names.zip"
local_filename = "babynames.zip"
if not os.path.exists(local_filename): # if the data exists don't download again
    with urllib.request.urlopen(data_url) as resp, open(local_filename, 'wb') as f:
        f.write(resp.read())
```

```
[18]: # Load data without unzipping the file
import zipfile
babynames = []
with zipfile.ZipFile(local_filename, "r") as zf:
    data_files = [f for f in zf.filelist if f.filename[-3:] == "txt"]
    def extract_year_from_filename(fn):
        return int(fn[3:7])
    for f in data_files:
        year = extract_year_from_filename(f.filename)
        with zf.open(f) as fp:
            df = pd.read_csv(fp, names=["Name", "Sex", "Count"])
            df["Year"] = year
            babynames.append(df)
babynames = pd.concat(babynames)

babynames.head()
```

```
[18]:
```

	Name	Sex	Count	Year
0	Mary	F	7065	1880
1	Anna	F	2604	1880
2	Emma	F	2003	1880

3	Elizabeth	F	1939	1880
4	Minnie	F	1746	1880

A bit of data cleaning

```
[19]: babynames['Name'] = babynames['Name'].str.lower()
babynames.tail()
```

```
[19]:
```

	Name	Sex	Count	Year
31949	zyheem	M	5	2019
31950	zykel	M	5	2019
31951	zyking	M	5	2019
31952	zyn	M	5	2019
31953	zyran	M	5	2019

How many people does this data represent?

```
[20]: format(babynames['Count'].sum(), ',d')
```

```
[20]: '355,149,899'
```

```
[21]: len(babynames)
```

```
[21]: 1989401
```

Is this number low or high?

It seems low. However the social security website states:

All names are from Social Security card applications for births that occurred in the United States after 1879. **Note that many people born before 1937 never applied for a Social Security card, so their names are not included in our data.** For others who did apply, our records may not show the place of birth, and again their names are not included in our data. All data are from a 100% sample of our records on Social Security card applications as of the end of February 2016.

Let's query to find rows that match desired conditions.

```
[22]: babynames[(babynames['Name'] == 'vela') & (babynames['Sex'] == 'F')].tail(5)
```

```
[22]:
```

	Name	Sex	Count	Year
6013	vela	F	22	2015
5594	vela	F	24	2016
7405	vela	F	16	2017
6213	vela	F	20	2018
6638	vela	F	18	2019

```
[23]: babynames[(babynames['Name'] == 'anthony') & (babynames['Year'] == 2000)]
```

```
[23]:
```

	Name	Sex	Count	Year
2782	anthony	F	52	2000
17673	anthony	M	19652	2000

```
[25]: babynames.query('Name.str.contains("data")', engine='python')
```

```
[25]:
```

	Name	Sex	Count	Year
9762	kidata	F	5	1975
24914	datavion	M	5	1995
23610	datavious	M	7	1997
12102	datavia	F	7	2000
27507	datavion	M	6	2001
28910	datari	M	5	2001
29138	datavian	M	5	2002
29139	datavious	M	5	2002
30572	datavion	M	5	2004
17139	datavia	F	5	2005
31027	datavion	M	5	2005
31021	datavion	M	6	2006
33338	datavious	M	5	2007
33339	datavius	M	5	2007
33402	datavious	M	5	2008
33081	datavion	M	5	2009
32497	datavious	M	5	2010

Subquestion: What is the proportion of Male and Female Individuals Over Time? In this example we construct a pivot table which aggregates the number of babies registered for each year by Sex.


```
[26]: pivot_year_name_count = pd.pivot_table(babynames,
        index=['Year'], # the row index
        columns=['Sex'], # the column values
        values='Count', # the field(s) to processed in each group
        aggfunc=np.sum,
    )

pivot_year_name_count.head()
```


```
[26]:
```

Sex	F	M
Year		
1880	90994	110490
1881	91953	100743
1882	107847	113686
1883	112319	104625
1884	129019	114442

```
[27]: pivot_year_name_count.plot(title='Names Registered that Year');
```




```
[28]: fig = go.Figure()
fig.add_trace(go.Scatter(x = pivot_year_name_count.index, y = pivot_year_name_count['F'], name = 'F', line=dict(color='gold')))
fig.add_trace(go.Scatter(x = pivot_year_name_count.index, y = pivot_year_name_count['M'], name = 'M', line=dict(color='blue')))
fig.update_layout(xaxis_title = 'Year', yaxis_title = 'Names Registered')
```



3.3.2 Subquestion: How many unique names for each year?

```
[29]: pivot_year_name_nunique = pd.pivot_table(babynames,
        index=['Year'],
        columns=['Sex'],
        values='Name',
        aggfunc=lambda x: len(np.unique(x)),
    )

pivot_year_name_nunique.plot(
    title='Number of Unique Names');
```



Some observations: - Registration data seems limited in the early 1900s. Because many people did not register before 1937. - You can see the baby boomers and the echo boom. - Females have greater diversity of names.

3.3.3 Subquestion: Computing the Proportion of Female Babies For Each Name:

```
[30]: sex_counts = pd.pivot_table(babynames, index='Name', columns='Sex',
        values='Count',
        aggfunc='sum', fill_value=0., margins=True)
sex_counts.head()
```

```
[30]: Sex      F      M  All
      Name
      aaban      0  120  120
      aabha     40    0   40
      aabid      0   16   16
      aabidah    5    0    5
      aabir      0   10   10
```

Compute proportion of female babies given each name.

```
[31]: prop_female = sex_counts['F'] / sex_counts['All']
      prop_female.head(10)
```

```
[31]: Name
      aaban      0.0
      aabha      1.0
      aabid      0.0
      aabidah    1.0
      aabir      0.0
      aabriella  1.0
      aada       1.0
      aadam      0.0
      aadan      0.0
      aadarsh    0.0
      dtype: float64
```

```
[32]: prop_female.tail(10)
```

```
[32]: Name
      zytavion    0.000000
      zytavious  0.000000
      zyus        0.000000
      zyva        1.000000
      zyvion      0.000000
      zyvon       0.000000
      zyyanna     1.000000
      zyyon       0.000000
      zzyzx       0.000000
      All        0.494913
      dtype: float64
```

Testing a few names

```
[33]: prop_female['audi']
```

```
[33]: 0.5978260869565217
```

```
[34]: prop_female['anthony']
```

```
[34]: 0.004856689867035234
```

```
[35]: prop_female['joey']
```

```
[35]: 0.1133165658350894
```

```
[36]: prop_female['mark']
```

```
[36]: 0.003307100877990732
```

```
[37]: prop_female["sarah"]
```

```
[37]: 0.9969322438050136
```

```
[38]: prop_female["min"]
```

```
[38]: 0.37598736176935227
```

```
[39]: prop_female["pat"]
```

```
[39]: 0.600140600694029
```

3.3.4 Modelling

Idea: Build Simple Classifier (Model) based on lookup table.

We can define a function to return the most likely Sex for a name. If there is an exact tie, the function returns Male. If the name does not appear in the social security dataset, we return Unknown.

```
[40]: def sex_from_name(name):  
      lower_name = name.lower()  
      if lower_name in prop_female.index:  
          return 'F' if prop_female[lower_name] > 0.5 else 'M'  
      else:  
          return "Unknown"
```

```
[41]: sex_from_name("audi")
```

```
[41]: 'F'
```

```
[42]: sex_from_name("joey")
```

```
[42]: 'M'
```

What fraction of students in IDSS this semester have names in the SSN dataset?


```
[43]: student_names = pd.Index(data["Firstname"]).intersection(prop_female.index)
print("Fraction of names in the babynames data:" , len(student_names) /
↪ len(data))
```

Fraction of names in the babynames data: 0.3655913978494624

Which names are not in the dataset? Why might these names not appear?

```
[44]: missing_names = pd.Index(data["Firstname"]).difference(prop_female.index)
print(missing_names)
print(len(missing_names))
```

```
Index(['*', '-', 'aikaterini', 'anhua', 'avinab', 'botao', 'boyang', 'boyuan',
      'buruo', 'ceyu',
      ...
      'zhehao', 'zhenyu', 'zhihan', 'zhoutian', 'zhu', 'zhuanghai', 'zhuo',
      'zidi', 'zijia', 'zixia'],
      dtype='object', length=205)
```


205

Observation: - That seems like a lot of missing names! - Should we continue with this dataset or find a better data source?

3.4 Q: What is the fraction of female and male students?

Back to the original and final question?

```
[45]: data['Pred. Sex'] = data['Firstname'].apply(sex_from_name)
(data[data['Pred. Sex'] != "Unknown"]['Pred. Sex'].value_counts()/
↪ len(data[data['Pred. Sex'] != "Unknown"])).plot(kind="barh");
```



3.5 6) Evaluation

- *Is this an accurate result ? Let's do an post-hoc survey and find out!*

Additionally: - Which visualisations do you need to communicate your findings? - How can you concisely describe the data, process, results and conclusions. - How can you summarise and document the steps you have taken? - What are the limitations of the analysis? ... should we rerun the whole analysis with a different strategy and perhaps data datasets?